

**Name: Husain Hashemi
roll no.: 20222727**

B.voc. Software Development - semester 7

**Data Analysis & Visualization
Practical File (1-8)**

Submitted to: Ms. Sonia Sodhi

Ramanujan College (DU) – November 2025

Q1

1. Write programs in Python using NumPy library to do the following:
 - a. Compute the mean, standard deviation, and variance of a two dimensional random integer array along the second axis.
 - b. Create a 2-dimensional array of size $m \times n$ integer elements, also print the shape, type and data type of the array and then reshape it into an $n \times m$ array, where n and m are user inputs given at the run time.
 - c. Test whether the elements of a given 1D array are zero, non-zero and NaN. Record the indices of these elements in three separate arrays.
 - d. Create three random arrays of the same size: Array1, Array2 and Array3. Subtract Array 2 from Array3 and store in Array4. Create another array Array5 having two times the values in Array1. Find Co-variance and Correlation of Array1 with Array4 and Array5 respectively.
 - e. Create two random arrays of the same size 10: Array1, and Array2. Find the sum of the first half of both the arrays and product of the second half of both the arrays.

Code a)

```
ria = np.random.randint(0,100, size=(4,5))
# print(ria)
print(np.mean(ria, axis=1))
print(np.var(ria, axis=1))
print(np.std(ria, axis=1))
```

Output

```
[55.6 55.4 60.8 30. ]
[311.44 624.24 438.56 316.4 ]
[17.64766273 24.98479538 20.94182418 17.78763616]
```

b)

```
n = int(input("give num for n cols: "))
m = int(input("give num for m rows: "))
d2 = np.array([[1,2,3,4,5],
               [6,7,8,9,10]])
print(d2.shape)
print(d2.ndim)
print(d2.dtype)
d2 = d2.reshape(m,1,n)
print(d2)
```

output

```
give num for n cols: 2
give num for m rows: 5
(2, 5)
2
int64
[[ 1  2]
 [ 3  4]
 [ 5  6]
 [ 7  8]
 [ 9 10]]
```

c)

```
d1 = np.array([1,2,3,4,5,0,6])
allzero = d1[d1 <= 0]
allnonzero = d1[d1 > 0]
nAn = d1[d1 == np.nan]
print(allzero)
print(allnonzero)
print(nAn)
```

Output

```
[0]
[1 2 3 4 5 6]
[]
```

d)

```
array1 = np.array([[1,2,3],[4,5,6]])
array2 = np.array([[7,8,9],[6,5,4]])
array3 = np.array([[5,6,8],[4,7,9]])

array4 = array2 - array3
array5 = array1 * 2
print(array4)
print(array5)
print(np.cov(array1,array4))
print(np.corrcoef(array1,array4))
print(np.cov(array1,array5))
print(np.corrcoef(array1,array5))
```

Output

```
[[ 2  2  1]
 [ 2 -2 -5]]
[[ 2  4  6]
 [ 8 10 12]]
[[ 1.          1.          -0.5         -3.5        ]
 [ 1.          1.          -0.5         -3.5        ]
 [-0.5         -0.5         0.33333333  1.66666667]
 [-3.5         -3.5         1.66666667 12.33333333]]
[[ 1.          1.          -0.8660254  -0.9966159 ]
 [ 1.          1.          -0.8660254  -0.9966159 ]
 [-0.8660254 -0.8660254   1.          0.82199494]
 [-0.9966159 -0.9966159   0.82199494  1.          ]]
[[1.  1.  2.  2.]
 [1.  1.  2.  2.]
 [2.  2.  4.  4.]
 [2.  2.  4.  4.]]
[[1.  1.  1.  1.]
 [1.  1.  1.  1.]
 [1.  1.  1.  1.]
 [1.  1.  1.  1.]]
```

e)

```
a1 = np.array([1,2,3,4,5,6,7,8,9,10])
a2 = np.array([10,9,8,7,6,5,4,3,2,1])
print(a1[0:5] + a2[0:5])
print(a1[5:] * a2[5:])
```

Output

```
[11 11 11 11 11]
[30 28 24 18 10]
```

Q2

Consider the following data frame containing a family name, gender of the family member and her/his monthly income in each record.

| Name | Gender | MonthlyIncome (Rs.) |
|-------|--------|---------------------|
| Shah | Male | 114000.00 |
| Vats | Male | 65000.00 |
| Vats | Female | 43150.00 |
| Kumar | Female | 69500.00 |
| Vats | Female | 155000.00 |
| Kumar | Male | 103000.00 |
| Shah | Male | 55000.00 |
| Shah | Female | 112400.00 |
| Kumar | Female | 81030.00 |
| Vats | Male | 71900.00 |

Write a program in Python using Pandas to perform the following:

- Calculate and display familywise gross monthly income.
- Display the highest and lowest monthly income for each family name
- Calculate and display monthly income of all members earning income less than Rs. 80000.00.
- Display total number of females along with their average monthly income
- Delete rows with Monthly income less than the average income of all members

a)

```
import pandas as pd

df = pd.read_csv("dff.csv")
print(df)

group_name = df.groupby("Name")
print(group_name["MonthlyIncome (Rs.)"].sum())
```

Output

```
Name  Gender  MonthlyIncome (Rs.)
0    Shah    Male        114000
1    Vats    Male        65000
2    Vats    Female     43150
3   Kumar   Female     69500
4    Vats    Female     155000
5   Kumar   Male        103000
6    Shah    Male        55000
7    Shah    Female     112400
8   Kumar   Female     81030
9    Vats    Male        71900
Name
Kumar    253530
Shah     281400
Vats     335050
Name: MonthlyIncome (Rs.), dtype: int64
```

b)

```
print("lowest income is: \n",group_name["MonthlyIncome (Rs.)"].min())
print("highest income is: \n",group_name["MonthlyIncome (Rs.)"].max())
```

Output

```
lowest income is:
Name
Kumar    69500
Shah     55000
Vats     43150
Name: MonthlyIncome (Rs.), dtype: int64
highest income is:
Name
Kumar    103000
Shah     114000
Vats     155000
Name: MonthlyIncome (Rs.), dtype: int64
```

c)

```
low_income = df[df["MonthlyIncome (Rs.)"] < 80000]
print(low_income)
```

Output

| | Name | Gender | MonthlyIncome (Rs.) |
|---|-------|--------|---------------------|
| 1 | Vats | Male | 65000 |
| 2 | Vats | Female | 43150 |
| 3 | Kumar | Female | 69500 |
| 6 | Shah | Male | 55000 |
| 9 | Vats | Male | 71900 |

d)

```
female = df[df['Gender'] == "Female"]
fe_sum = female["MonthlyIncome (Rs.)"].sum()
print(female)
print("all females tot income is: ",fe_sum)
```

Output

| | Name | Gender | MonthlyIncome (Rs.) |
|---|-------|--------|---------------------|
| 2 | Vats | Female | 43150 |
| 3 | Kumar | Female | 69500 |
| 4 | Vats | Female | 155000 |
| 7 | Shah | Female | 112400 |
| 8 | Kumar | Female | 81030 |

all females tot income is: 461080

e)

```
avg_income = df["MonthlyIncome (Rs.)"].mean()
print(avg_income)
df_filter = df[df["MonthlyIncome (Rs.)"] > avg_income]
print(df_filter)
```

Output:

| | Name | Gender | MonthlyIncome (Rs.) |
|---|-------|--------|---------------------|
| 0 | Shah | Male | 114000 |
| 4 | Vats | Female | 155000 |
| 5 | Kumar | Male | 103000 |
| 7 | Shah | Female | 112400 |

Q3

Use a dataset of your choice from Open Data Portal (<https://data.gov.in/>, UCI repository) or load from scikit, seaborn library for the following exercises to practice the concepts learnt.

- Load a Pandas dataframe with a selected dataset. Identify and count the missing values in a dataframe. Clean the data after removing noise as follows
 - a) Drop duplicate rows.
 - b) Detect the outliers and remove the rows having more than two outliers identified using boxplot.
 - c) Identify the most correlated positively correlated attributes and negatively correlated attributes

Code:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df = sns.load_dataset("titanic")
print(df.head())
```

I: Identify and count missing values

```
missing_vals = df.isnull().sum()
print(missing_vals)
```

II: Drop duplicate rows

```
dropped_dup = df.drop_duplicates()
print("\nDataFrame shape after dropping duplicates:", dropped_dup.shape)
```

Step 3: Detect outliers using boxplot

```
num_cols = df_clean.select_dtypes(include=[np.number]).columns
print("\nNumerical columns considered for outlier detection:", list(num_cols))

# Detect outliers based on IQR for each numerical column
def detect_outliers_iqr(data, col):
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return (data[col] < lower_bound) | (data[col] > upper_bound)
```

```
# Create a DataFrame to hold outlier flags for each numerical column
outlier_flags = pd.DataFrame()
for col in num_cols:
    outlier_flags[col] = detect_outliers_iqr(df_clean, col)
```

```
# Count outliers per row
outlier_flags['outlier_count'] = outlier_flags.sum(axis=1)
```

```
# Remove rows having more than two outliers
df_no_outliers = df_clean[outlier_flags['outlier_count'] <= 2]
print("\nDataFrame shape after removing rows with more than two outliers:", df_no_outliers.shape)
```

```
# Optional: Plot boxplots before and after cleaning (you can skip this)
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
sns.boxplot(x=df_clean['age'])
plt.title("Age - Before Outlier Removal")

plt.subplot(1, 2, 2)
sns.boxplot(x=df_no_outliers['age'])
plt.title("Age - After Outlier Removal")
plt.show()
```

Step 4: Correlation matrix and find most correlated positive and negative pairs

```
# Step 4: Correlation matrix and find most correlated positive and negative pairs
corr_matrix = df_no_outliers[num_cols].corr()
print("\nCorrelation matrix:\n", corr_matrix)
```

```
# Extract upper triangle of the correlation matrix without diagonal
upper_tri = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
# Find the most positively correlated pair
max_corr = upper_tri.max().max()
max_pair = [(col, row) for col in upper_tri.columns for row in upper_tri.index if upper_tri.loc[row, col] == max_corr]

# Find the most negatively correlated pair
min_corr = upper_tri.min().min()
min_pair = [(col, row) for col in upper_tri.columns for row in upper_tri.index if upper_tri.loc[row, col] == min_corr]
print(f"\nMost positively correlated attributes: {max_pair} with correlation {max_corr:.3f}")
print(f"Most negatively correlated attributes: {min_pair} with correlation {min_corr:.3f}")
```

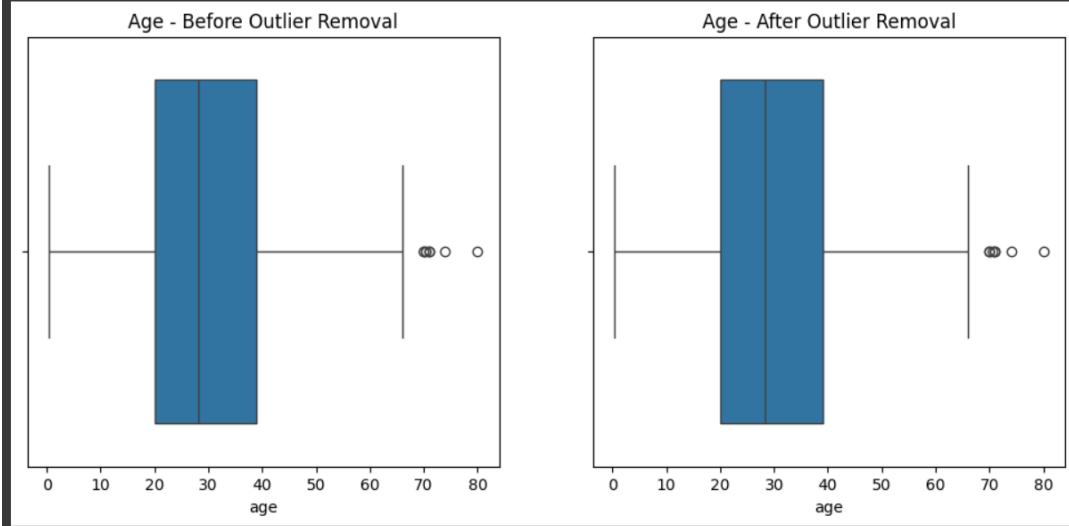
output:

```
sibsp      0  
parch     0  
fare      0  
embarked   2  
class      0  
who        0  
adult_male  0  
deck       688  
embark_town 2  
alive      0  
alone      0  
dtype: int64
```

DataFrame shape after dropping duplicates: (784, 15)

Numerical columns considered for outlier detection: ['survived', 'pclass', 'age', 'sibsp', 'parch', 'fare']

DataFrame shape after removing rows with more than two outliers: (784, 15)



Correlation matrix:

| | survived | pclass | age | sibsp | parch | fare |
|----------|-----------|-----------|-----------|-----------|-----------|-----------|
| survived | 1.000000 | -0.332658 | -0.086548 | -0.036589 | 0.070307 | 0.246769 |
| pclass | -0.332658 | 1.000000 | -0.369361 | 0.088014 | 0.040296 | -0.549216 |
| age | -0.086548 | -0.369361 | 1.000000 | -0.315116 | -0.195036 | 0.092707 |
| sibsp | -0.036589 | 0.088014 | -0.315116 | 1.000000 | 0.381433 | 0.135147 |
| parch | 0.070307 | 0.040296 | -0.195036 | 0.381433 | 1.000000 | 0.191942 |
| fare | 0.246769 | -0.549216 | 0.092707 | 0.135147 | 0.191942 | 1.000000 |

Most positively correlated attributes: [('parch', 'sibsp')] with correlation 0.381

Most negatively correlated attributes: [('fare', 'pclass')] with correlation -0.549

Q4

- Import iris data using sklearn library or (Download IRIS data from:
<https://archive.ics.uci.edu/ml/datasets/iris> or import it from sklearn.datasets)
- Compute mean, mode, median, standard deviation, confidence interval and standard error for each feature
 - Compute correlation coefficients between each pair of features and plot heatmap
 - Find covariance between length of sepal and petal iv. Build contingency table for class feature

Code

```
import pandas as pd
import numpy as np
from sklearn import datasets
from scipy import stats
import matplotlib.pyplot as plt
import seaborn as sns

iris = datasets.load_iris()
df = pd.DataFrame(data = iris.data, columns=iris.feature_names)
print(df.head())
```

```
df['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)
```

(a) Descriptive Statistics

```
# Mean
mean_vals = df.iloc[:, :].mean()
print(mean_vals)
```

```
# # Mode
mode_vals = df.iloc[:, :].mean().iloc[0]
print(mode_vals)
```

```
# Median
med_vals = df.iloc[:, :].median()
print(med_vals)
```

```
# Standard variation
std_vals = df.iloc[:, :].std()
print(std_vals)
```

```
# Standard error
st_error = df.iloc[:, :].sem()
print(st_error)
```

```

# 95% Confidence Interval
conf_int = {}
for col in df.columns[:-1]:
    mean = df[col].mean()
    se = df[col].sem()
    ci = stats.t.interval(0.95, len(df[col])-1, loc=mean, scale=se)
    conf_int[col] = ci
print("\n95% Confidence Intervals:")
for col, ci in conf_int.items():
    print(f"{col}: {ci}")

```

(b) Correlation Coefficients + Heatmap

```

# (b) Correlation Coefficients + Heatmap
print("\n(b) Correlation Matrix:")

corr_matrix = df.iloc[:, :].corr()
print(corr_matrix)

# Plot heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Heatmap of Feature Correlations")
plt.show()

```

(c) Covariance: Sepal Length vs Petal Length

```

# (c) Covariance: Sepal Length vs Petal Length
print("\n(c) Covariance between Sepal Length and Petal Length:")

cov_matrix = np.cov(df["sepal length (cm)"], df["petal length (cm)"])
cov_val = cov_matrix[0, 1]
print("Covariance:", cov_val)

```

(d) Contingency Table for Class Feature

```

# (d) Contingency Table for Class Feature
print("\n(d) Contingency Table:")

contingency_table = pd.crosstab(index=df['species'], columns="count")
print(contingency_table)

```

Output:

```
      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0                  5.1              3.5                 1.4               0.2
1                  4.9              3.0                 1.4               0.2
2                  4.7              3.2                 1.3               0.2
3                  4.6              3.1                 1.5               0.2
4                  5.0              3.6                 1.4               0.2
PS C:\Users\Husain\Desktop\gg> []
```

(a) Descriptive Statistics:

Mean:

```
sepal length (cm)      5.843333
sepal width (cm)       3.057333
petal length (cm)      3.758000
petal width (cm)       1.199333
dtype: float64
```

Mode:

```
sepal length (cm)      5.0
sepal width (cm)       3.0
petal length (cm)      1.4
petal width (cm)       0.2
Name: 0, dtype: float64
```

Median:

```
sepal length (cm)      5.80
sepal width (cm)       3.00
petal length (cm)      4.35
petal width (cm)       1.30
dtype: float64
```

Standard Deviation:

```
sepal length (cm)      0.828066
sepal width (cm)       0.435866
petal length (cm)      1.765298
petal width (cm)       0.762238
dtype: float64
```

Standard Error:

```
sepal length (cm)      0.067611
sepal width (cm)       0.035588
petal length (cm)      0.144136
petal width (cm)       0.062236
dtype: float64
```

95% Confidence Intervals:

```
sepal length (cm): (np.float64(5.709732481507366), np.float64(5.976934185159301))
sepal width (cm): (np.float64(2.9870103180785432), np.float64(3.127656348588124))
petal length (cm): (np.float64(3.473185370199511), np.float64(4.04281462980049))
petal width (cm): (np.float64(1.0763532977706853), np.float64(1.3223133688959818))
```

(b) Correlation Matrix:

```
      sepal length (cm)  sepal width (cm)  petal length (cm) \
sepal length (cm)          1.000000      -0.117570      0.871754
sepal width (cm)         -0.117570       1.000000     -0.428440
petal length (cm)         0.871754      -0.428440      1.000000
petal width (cm)          0.817941      -0.366126      0.962865

                           petal width (cm)
sepal length (cm)          0.817941
sepal width (cm)         -0.366126
petal length (cm)          0.962865
petal width (cm)          1.000000
```



(c) Covariance between Sepal Length and Petal Length:
 Covariance: 1.2743154362416111

(d) Contingency Table:

| col_0 | count |
|------------|-------|
| species | |
| setosa | 50 |
| versicolor | 50 |
| virginica | 50 |

Q5

Consider two data files (in CSV format) having attendance of two workshops. Each file has three fields ‘Name’, ‘Date, duration (in minutes) where names are unique within a file. Note that duration may take one of three values (30, 40, 50) only. Import the data into two data frames and do the following:

- a. Perform merging of the two data frames to find the names of students who had attended both workshops.
- b. Find names of all students who have attended a single workshop only.
- c. Merge two data frames row-wise and find the total number of records in the data frame.
- d. Merge two data frames row-wise and use two columns viz. names and dates as multi-row indexes. Generate descriptive statistics for this hierarchical data frame.

Code:

```
import pandas as pd

df1 = pd.read_csv("workshop1.csv")
df2 = pd.read_csv("workshop2.csv")

print("Workshop 1 Data:\n", df1, "\n")
print("Workshop 2 Data:\n", df2, "\n")
```

(a) Find names of students who attended BOTH workshops

```
# (a) Find names of students who attended BOTH workshops
both = pd.merge(df1, df2, on="Name", suffixes=("_w1", "_w2"))
print("(a) Students who attended both workshops:\n", both["Name"].tolist(), "\n")
```

(b) Find names of students who attended ONLY ONE workshop

```
# (b) Find names of students who attended ONLY ONE workshop
# Use outer merge and find where Name appears only once
merged_outer = pd.merge(df1, df2, on="Name", how="outer", indicator=True)
single_only = merged_outer[merged_outer["_merge"] != "both"]["Name"]
print("(b) Students who attended a single workshop only:\n", single_only.tolist(), "\n")
```

(c) Merge two data frames row-wise (stack them)

```
# (c) Merge two data frames row-wise (stack them)
merged_rowwise = pd.concat([df1, df2], axis=0)
print("(c) Row-wise merged DataFrame:\n", merged_rowwise, "\n")
print("Total number of records:", len(merged_rowwise), "\n")
```

```
# (d) Merge row-wise again, set Name and Date as multi-row index
```

```
# (d) Merge row-wise again, set Name and Date as multi-row index
#      and generate descriptive statistics
df_hierarchical = merged_rowwise.set_index(["Name", "Date"])
print("(d) Hierarchical DataFrame:\n", df_hierarchical, "\n")

print("Descriptive Statistics for Hierarchical DataFrame:")
print(df_hierarchical.describe())
```

Output:

```
Workshop 1 Data:
      Name      Date Duration
0    Rudra  08-10-2025      40
1   Husain  08-10-2025      38
2     Amit  08-10-2025      40
3  Mozammil  08-10-2025      30
4  Rudransh  08-10-2025      25
```

```
Workshop 2 Data:
      Name      Date Duration
0   Ankush  09-10-2025      50
1   Pankaj  09-10-2025      45
2    Dhruv  09-10-2025      40
3   Akshat  09-10-2025      30
4    Rudra  09-10-2025      35
```

(a) Students who attended both workshops:
['Rudra']

(b) Students who attended a single workshop only:
['Akshat', 'Amit', 'Ankush', 'Dhruv', 'Husain', 'Mozammil', 'Pankaj', 'Rudransh']

(c) Row-wise merged DataFrame:

| | Name | Date | Duration |
|---|----------|------------|----------|
| 0 | Rudra | 08-10-2025 | 40 |
| 1 | Husain | 08-10-2025 | 38 |
| 2 | Amit | 08-10-2025 | 40 |
| 3 | Mozammil | 08-10-2025 | 30 |
| 4 | Rudransh | 08-10-2025 | 25 |
| 0 | Ankush | 09-10-2025 | 50 |
| 1 | Pankaj | 09-10-2025 | 45 |
| 2 | Dhruv | 09-10-2025 | 40 |
| 3 | Akshat | 09-10-2025 | 30 |
| 4 | Rudra | 09-10-2025 | 35 |

Total number of records: 10

(d) Hierarchical DataFrame:

| Name | Date | Duration |
|----------|------------|----------|
| Rudra | 08-10-2025 | 40 |
| Husain | 08-10-2025 | 38 |
| Amit | 08-10-2025 | 40 |
| Mozammil | 08-10-2025 | 30 |
| Rudransh | 08-10-2025 | 25 |
| Ankush | 09-10-2025 | 50 |
| Pankaj | 09-10-2025 | 45 |
| Dhruv | 09-10-2025 | 40 |
| Akshat | 09-10-2025 | 30 |
| Rudra | 09-10-2025 | 35 |

Descriptive Statistics for Hierarchical DataFrame:

| | Duration |
|-------|-----------|
| count | 10.000000 |
| mean | 37.300000 |
| std | 7.498889 |
| min | 25.000000 |
| 25% | 31.250000 |
| 50% | 39.000000 |
| 75% | 40.000000 |
| max | 50.000000 |

Q6

- Load Titanic data from sklearn library , plot the following with proper legend and axis labels:
 - Plot bar chart to show the frequency of survivors and non-survivors for male and female passengers separately
 - Draw a scatter plot for any two selected features
 - Compare density distribution for features age and passenger fare
 - Use a pair plot to show pairwise bivariate distribution

Code

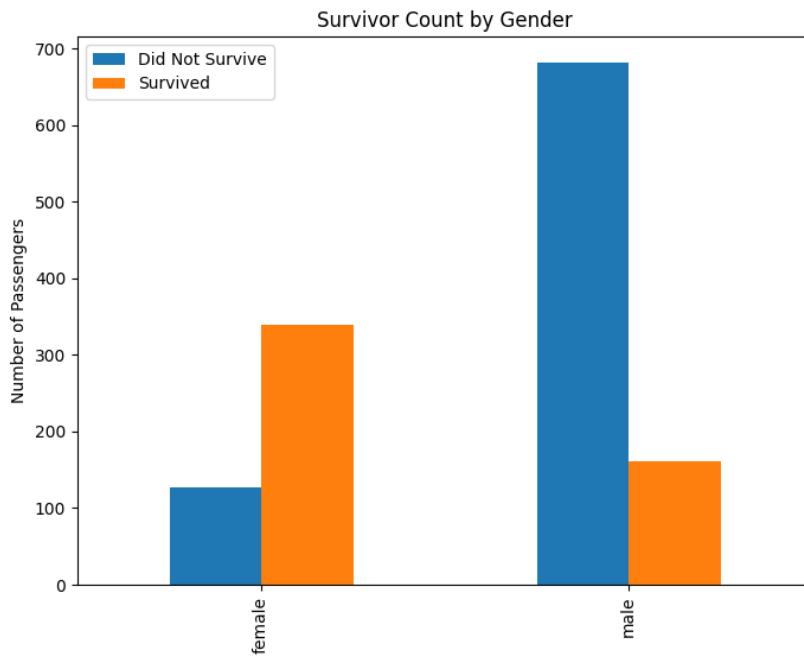
```
from sklearn.datasets import fetch_openml
import matplotlib.pyplot as plt
import seaborn as sns

titanic = fetch_openml('titanic', version=1, as_frame=True)
df = titanic.frame
```

A)

```
# Count survivors by gender
survivor_counts = df.groupby(['sex', 'survived']).size().unstack()
# Plot
survivor_counts.plot(kind='bar', figsize=(8,6))
plt.title('Survivor Count by Gender')
plt.xlabel('Gender')
plt.ylabel('Number of Passengers')
plt.legend(['Did Not Survive', 'Survived'])
plt.show()
```

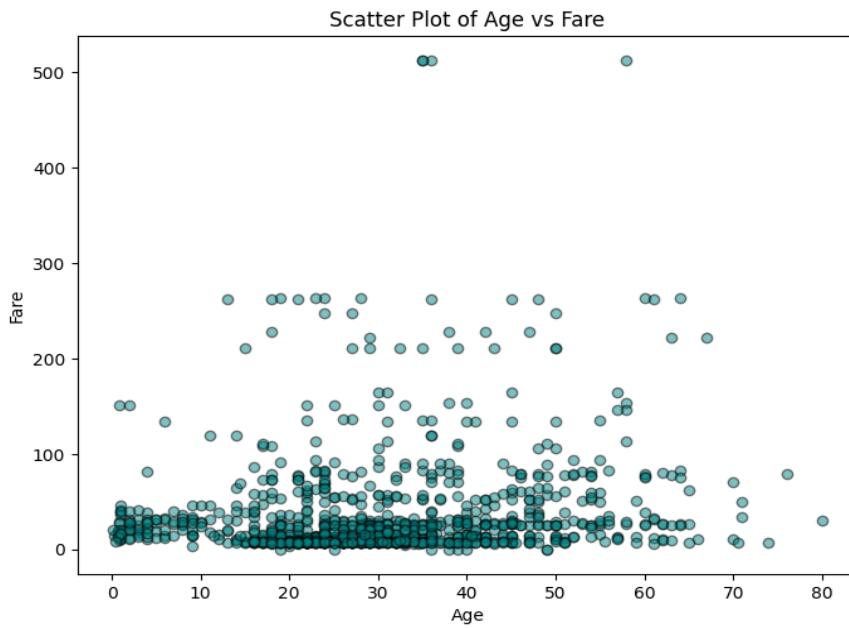
Output:



B)

```
# Drop missing values for scatter plot
df_scatter = df[['age', 'fare']].dropna()
plt.figure(figsize=(8,6))
plt.scatter(df_scatter['age'], df_scatter['fare'], alpha=0.5, c='teal', edgecolor='k')
plt.title('Scatter Plot of Age vs Fare')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.show()
```

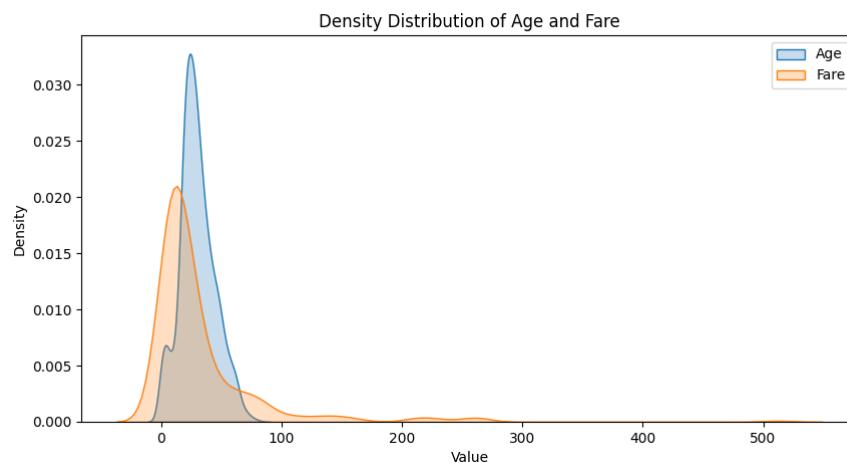
Output:



C)

```
plt.figure(figsize=(10,5))
sns.kdeplot(df['age'].dropna(), label='Age', fill=True)
sns.kdeplot(df['fare'].dropna(), label='Fare', fill=True)
plt.title('Density Distribution of Age and Fare')
plt.xlabel('Value')
plt.ylabel('Density')
plt.legend()
plt.show()
```

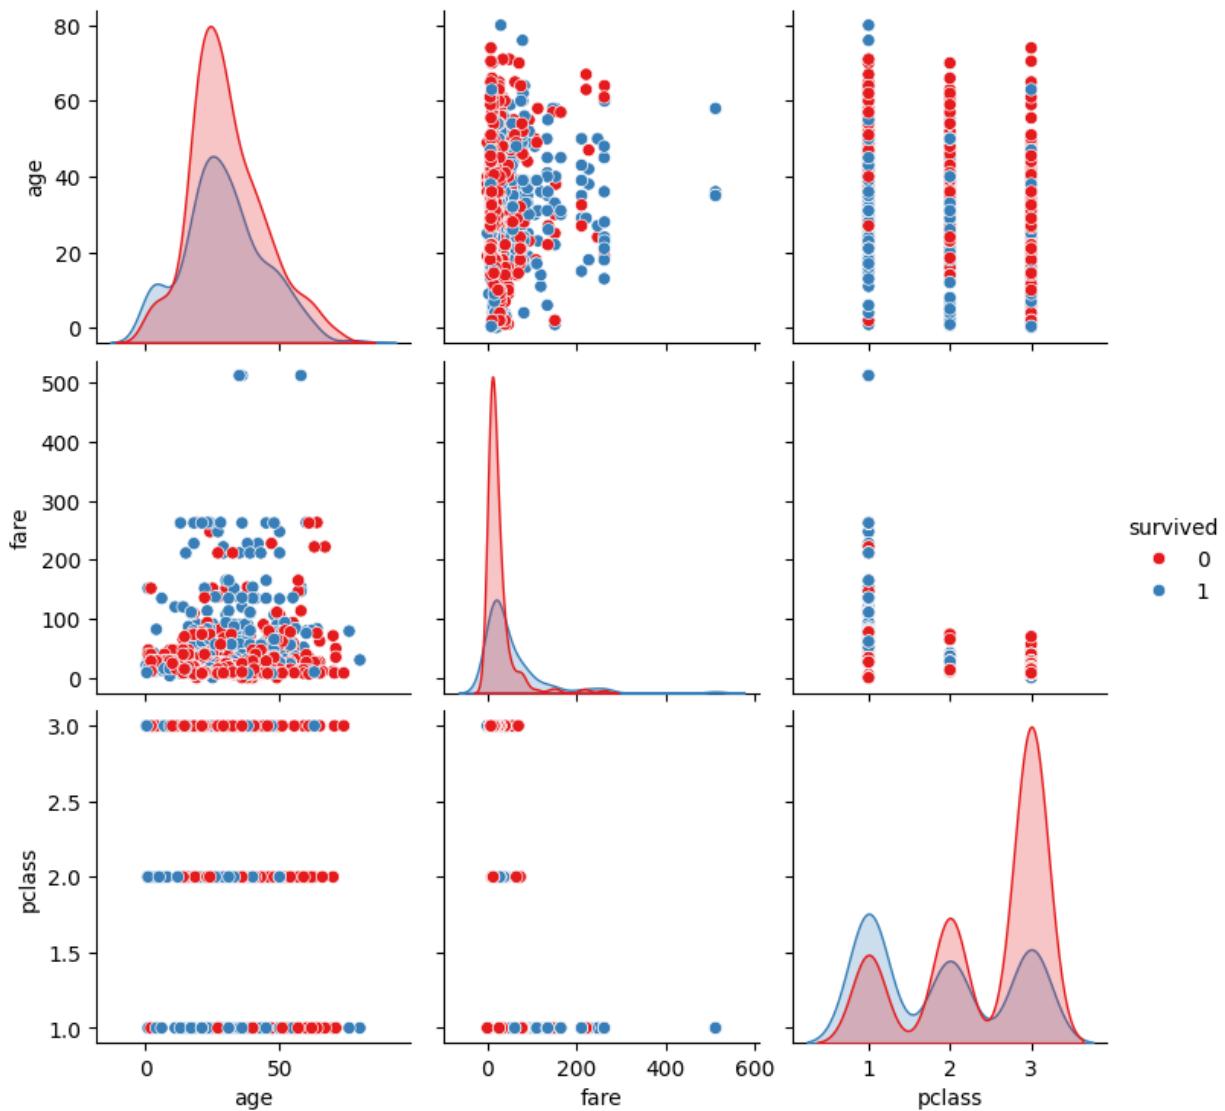
Output:



D)

```
numeric_features = ['age', 'fare', 'pclass', 'survived']
sns.pairplot(df[numeric_features].dropna(), hue='survived', palette='Set1', diag_kind='kde')
plt.suptitle('Pairwise Bivariate Distribution', y=1.02)
plt.show()
```

Output:



Q7

- Using Titanic dataset, do the following
- Find total number of passengers with age less than 30
 - Find total fare paid by passengers of first class
 - Compare number of survivors of each passenger class

Code:

```
from sklearn.datasets import fetch_openml
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

titanic = fetch_openml('titanic', version=1, as_frame=True)
df = titanic.frame
```

```
# Convert 'survived' column to numeric
df['survived'] = pd.to_numeric(df['survived'])

# Drop missing ages before filtering
num_passengers_under_30 = df[df['age'].notna() & (df['age'] < 30)].shape[0]
print("Total number of passengers with age less than 30:", num_passengers_under_30)
```

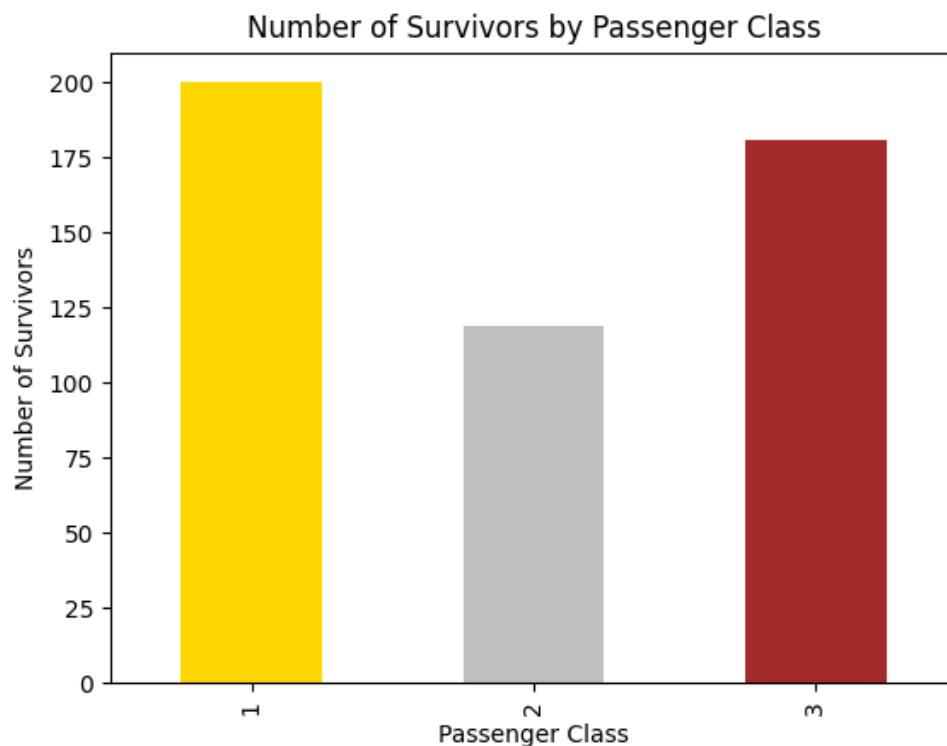
```
total_fare_first_class = df[df['pclass'] == 1]['fare'].sum()
print("Total fare paid by first-class passengers:", total_fare_first_class)

survivors_by_class = df[df['survived'] == 1].groupby('pclass').size()
print("Number of survivors by passenger class:")
print(survivors_by_class)
```

```
survivors_by_class.plot(kind='bar', color=['gold', 'silver', 'brown'])
plt.title('Number of Survivors by Passenger Class')
plt.xlabel('Passenger Class')
plt.ylabel('Number of Survivors')
plt.show()
```

Output:

```
Total number of passengers with age less than 30: 569
Total fare paid by first-class passengers: 28265.404300000002
Number of survivors by passenger class:
pclass
1    200
2    119
3    181
dtype: int64
```



Q8

5. Download any dataset and do the following
- Count number of categorical and numeric features
 - Remove one correlated attribute (if any)
 - Display five-number summary of each attribute and show it visually

Code:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Load dataset
df = sns.load_dataset('penguins')
```

```
# Identify numeric and categorical columns
numeric_features = df.select_dtypes(include=[np.number]).columns
categorical_features = df.select_dtypes(include=['object', 'category']).columns
print("Number of numeric features:", len(numeric_features))
print("Number of categorical features:", len(categorical_features))
print("Numeric features:", list(numeric_features))
print("categorical features:", list(categorical_features))
```

```
# Correlation matrix
corr_matrix = df[numeric_features].corr()
print("Correlation matrix:\n", corr_matrix)
```

```
# Example: if correlation > 0.8, drop one column
# Here, let's check for correlations > 0.8 (excluding self-correlation)
high_corr = np.where((corr_matrix.abs() > 0.8) & (corr_matrix.abs() < 1))
high_corr_pairs = [(numeric_features[i], numeric_features[j]) for i,j in zip(*high_corr) if i < j]
print("Highly correlated pairs:", high_corr_pairs)
# If any, remove the second column of the first pair
if high_corr_pairs:
    col_to_drop = high_corr_pairs[0][1]
    df = df.drop(columns=[col_to_drop])
    print(f"Dropped column '{col_to_drop}' due to high correlation")
    # Update numeric_features after dropping the column
    numeric_features = df.select_dtypes(include=[np.number]).columns
```

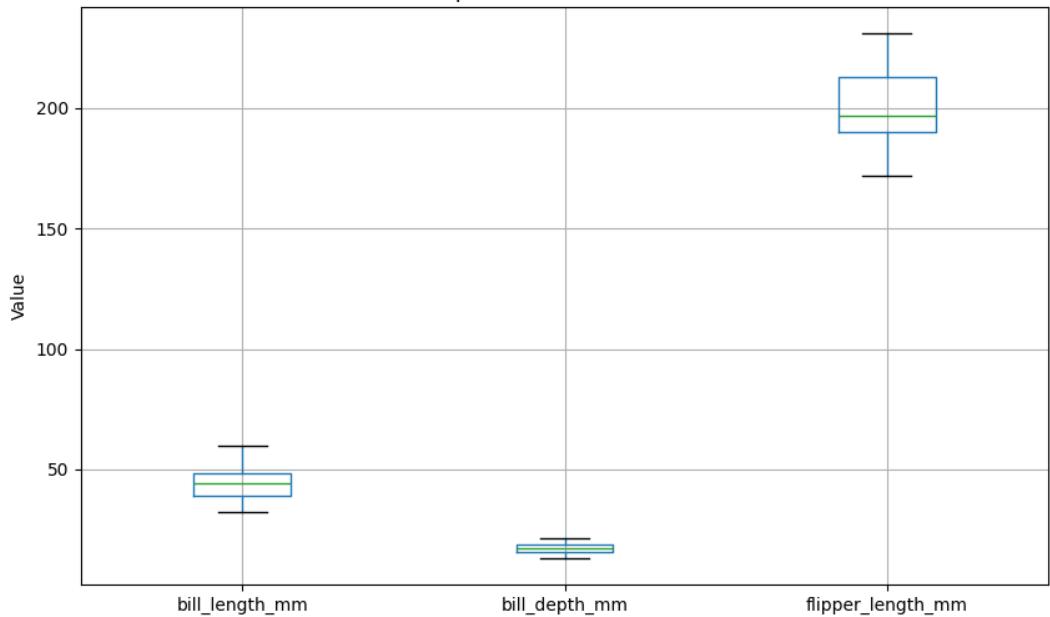
```
# Five-number summary
five_num_summary = df[numeric_features].describe().loc[['min','25%','50%','75%','max']]
print("Five-number summary:\n", five_num_summary)
```

```
# Visualize using boxplots
plt.figure(figsize=(10,6))
df[numeric_features].boxplot()
plt.title("Boxplot of Numeric Features")
plt.ylabel("Value")
plt.show()
df[numeric_features].hist(figsize=(12,6), bins=15)
plt.suptitle("Histogram of Numeric Features")
plt.show()
```

Output:

```
Number of numeric features: 4
Number of categorical features: 3
Numeric features: ['bill_length_mm', 'bill_depth_mm', 'flipper_length_mm', 'body_mass_g']
Categorical features: ['species', 'island', 'sex']
Correlation matrix:
      bill_length_mm  bill_depth_mm  flipper_length_mm  body_mass_g
bill_length_mm    1.000000   -0.235053     0.656181    0.595110
bill_depth_mm    -0.235053    1.000000    -0.583851   -0.471916
flipper_length_mm    0.656181   -0.583851     1.000000    0.871202
body_mass_g       0.595110   -0.471916     0.871202    1.000000
Highly correlated pairs: [('flipper_length_mm', 'body_mass_g')]
Dropped column 'body_mass_g' due to high correlation
Five-number summary:
      bill_length_mm  bill_depth_mm  flipper_length_mm
min        32.100        13.1          172.0
25%       39.225        15.6          190.0
50%       44.450        17.3          197.0
75%       48.500        18.7          213.0
max       59.600        21.5          231.0
```

Boxplot of Numeric Features



Histogram of Numeric Features

